

COMP9444 Neural Networks and Deep Learning

Session 2, 2018

Solutions to Exercises 1: Perceptrons


This page was last updated: 07/26/2018 11:37:03

1. Perceptron Learning

- a. Construct by hand a Perceptron which correctly classifies the following data; use your knowledge of plane geometry to choose appropriate values for the weights w_0 , w_1 and w_2 .

| Training Example | x_1 | x_2 | Class |
|------------------|-------|-------|-------|
| a. | 0 | 1 | -1 |
| b. | 2 | 0 | -1 |
| c. | 1 | 1 | +1 |

The first step is to plot the data on a 2-D graph, and draw a line which separates the positive from the negative data points:



This line has slope $-1/2$ and x_2 -intercept $5/4$, so its equation is:


$$x_2 = 5/4 - x_1/2,$$

i.e. $2x_1 + 4x_2 - 5 = 0$.

Taking account of which side is positive, this corresponds to these weights:

$$\begin{aligned} w_0 &= -5 \\ w_1 &= 2 \\ w_2 &= 4 \end{aligned}$$

Alternatively, we can derive weights $w_1=1$ and $w_2=2$ by drawing a vector normal to the separating line, in the direction pointing towards the positive data points:



The bias weight w_0 can then be found by computing the dot product of the normal vector with a perpendicular vector from the separating line to the origin. In this case $w_0 = 1(-0.5) + 2(-1) = -2.5$

(Note: these weights differ from the previous ones by a normalizing constant, which is fine for a Perceptron)

- b. Demonstrate the Perceptron Learning Algorithm on the above data, using a learning rate of 1.0 and initial weight values of

$$w_0 = -1.5$$

$$w_1 = 0$$

$$w_2 = 2$$

In your answer, you should clearly indicate the new weight values at the end of each training step.


| Iteration | w_0 | w_1 | w_2 | Training Example | x_1 | x_2 | Class | $s = w_0 + w_1x_1 + w_2x_2$ | Action |
|-----------|-------|-------|-------|------------------|-------|-------|-------|-----------------------------|----------|
| 1 | -1.5 | 0 | 2 | a. | 0 | 1 | - | +0.5 | Subtract |
| 2 | -2.5 | 0 | 1 | b. | 2 | 0 | - | -2.5 | None |
| 3 | -2.5 | 0 | 1 | c. | 1 | 1 | + | -1.5 | Add |
| 4 | -1.5 | 1 | 2 | a. | 0 | 1 | - | +0.5 | Subtract |
| 5 | -2.5 | 1 | 1 | b. | 2 | 0 | - | -0.5 | None |
| 6 | -2.5 | 1 | 1 | c. | 1 | 1 | + | -0.5 | Add |
| 7 | -1.5 | 2 | 2 | a. | 0 | 1 | - | +0.5 | Subtract |
| 8 | -2.5 | 2 | 1 | b. | 2 | 0 | - | +1.5 | Subtract |
| 9 | -3.5 | 0 | 1 | c. | 1 | 1 | + | -2.5 | Add |
| 10 | -2.5 | 1 | 2 | a. | 0 | 1 | - | -0.5 | None |
| 11 | -2.5 | 1 | 2 | b. | 2 | 0 | - | -0.5 | None |
| 12 | -2.5 | 1 | 2 | c. | 1 | 1 | + | +0.5 | None |

2. XOR Network

Construct by hand a Neural Network (or Multi-Layer Perceptron) that computes the XOR function of two inputs. Make sure the connections, weights and biases of your

network are clearly visible.


There are a number of ways to express XOR as a combination of simpler functions that are linearly separable. For example, using NOR as an abbreviation for "NOT OR", $(x_1 \text{ XOR } x_2)$ can be written as $(x_1 \text{ AND } x_2) \text{ NOR } (x_1 \text{ NOR } x_2)$. This decomposition allows us to compute XOR with a network like this:



Challenge: Can you construct a Neural Network to compute XOR which has only one hidden unit, but also includes shortcut connections from the two inputs directly to the (one) output.

Hint: start with a network that computes the inclusive OR, and then try to think of how it could be modified.

Exclusive OR (XOR) is very similar to normal (inclusive) OR, except for the case where both inputs are True, i.e. where $(x_1 \text{ AND } x_2)$ is True. We therefore introduce a single hidden unit which computes $(x_1 \text{ AND } x_2)$. This hidden unit is connected to the output with a negative weight, thus forcing that case to be classified negatively.



The addition of this hidden "feature" creates a 3-dimensional space in which the points can be linearly separated by a plane. The weights for the

output unit (+1,+1,-2) specify a vector perpendicular to the separating plane, and its distance from the origin is determined by the output bias divided by the length of this vector.

